

Guarded publishing

Let an agent draft in public, and make it impossible to say the wrong thing.

BUILT AND TESTED · NOTHING PUBLISHED YET

SCHEMATIC

BLOCKED

employer

platform

stakeholder

programme

ALLOWED

shipped a thing

what broke

what it cost

22 / 22 guard tests. A hard monthly ceiling, enforced in code.

22 / 22

GUARD REGRESSION TESTS

25

HARD MONTHLY CEILING

0

THIRD-PARTY DEPENDENCIES

1

TYPED APPROVAL PER POST

WHAT IT IS

The hard part of posting about your work is writing it, not clicking publish - so the drafting is automated and the publishing stays behind a typed human approval. Two rules are enforced in code rather than by good intentions: a word-boundary deny-list that blocks any employer term from ever leaving the machine, and a hard monthly ceiling that refuses to spend past it.

WHAT IT CHANGES

Confidential terms are blocked mechanically, not by remembering to check

A guard that cries wolf gets switched off, so false-positive tests are permanent

Spend is capped in code, so an agent cannot run up a bill

Python standard library · OAuth 1.0a · regression suite

Juan E. Diaz R. · Aruba-based · Remote worldwide · alphaowl01@gmail.com

SOMETHING LIKE THIS, FOR YOUR WORK?

The interesting part is your version of the problem.

Tell me the workflow that costs you the most time or the most trust. I will tell you honestly whether a system fixes it.

[Start a conversation](#)